

Predict Hacker News Upvotes

Feature Fusion

Predict Hacker News Upvotes

4. ▲ The Best Essay (paulgraham.com)
207 points by tosh 11 hours ago | hide | 143 comments

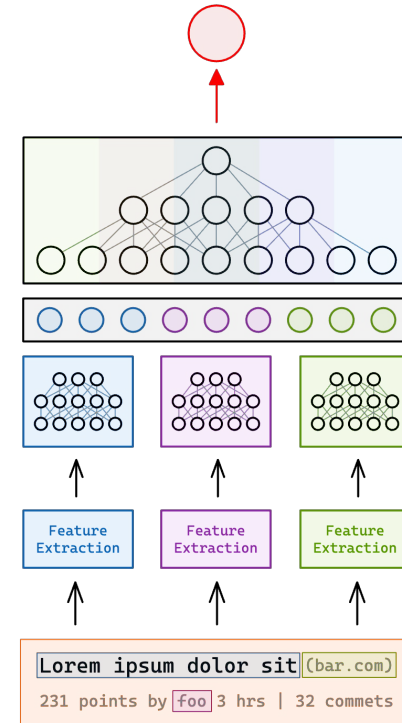


→ 265

Ubiquitous supervised learning task:

- Content Recommendation Systems
- Email & Message Prioritization
- Homework Assignment Score
- Spam Detection Systems
- ...

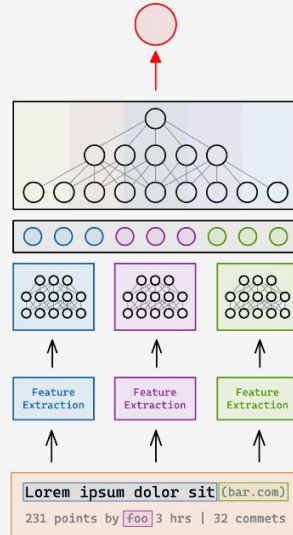
But also a Reward Model where on top we can train a Policy Model in a RLHF setup



Predict Hacker News Upvotes


Late Fusion

- Pro: Specific modality processing
- Pro: Independent specialised training
- Con: Requires separate data
- Con: May miss modalities interactions




Early Fusion


- Pro: Simple single train dataset
- Pro: Captures modality interactions
- Pro: Easy with SOTA models
- Con: Requires very complex models



Predict Hacker News Upvotes



CBOW



Skip-gram

Predict Hacker News Upvotes

Order of execution:

1. Exploratory Data Analysis (EDA)
 - a. look for cues
 - b. understand the distributions
 - c. data quality issues
2. Build prediction models
 - a. pre-train word2vec
 - b. finetune embeddings
 - c. regression task
3. Deploy a predictor service
 - a. grab a server on Computa
 - b. run Docker container
 - c. respond to live requests

Predict Hacker News Upvotes

Tools you can use:

- PyTorch, Python, FastAPI
- Bash, Docker
- Word2Vec

You can find the data:

- postgres://sy91dzb:g5t49ao@178.156.142.230:5432/hd64m1ki
 - **SELECT * FROM "hacker_news"."items" LIMIT 1;**
 - "hacker_news"."items" → 28GB of data, ~42 million rows
 - "hacker_news"."users" → ~15MB, eventually ~1 million

Predict Hacker News Upvotes



Word2Vec and Hacker
News Upvote
Prediction

4 Lessons | 40 hours
with **Ardavan Afshar**

Icons: Bar chart, Bookshelf, Gear, Refresh, Speaker

<https://cortex.mlx.institute/>

Good luck!

Recurrent Rebels

Anton Dergunov
Nikolas Kuhn
Tyrone Nicholas
Rasched Haidari

Perceptron Party

Jingyan Chen
Melanie Wong
Joao Esteves
James Carter

Kernel Kittens

Peter O'Keeffe
Ben Williams
Miguel Parracho
Oleh Khomiak

Gradient Gigglers

Helen Zhou
Prima Gouse
Andrei Zhirnov
Hikaru Tsujimura

Backprop Bunch

Kadriye Turkcan
Yali Pan
David Edey
Ben Liong

Bayesian Buccaneers

Tomas Krajcoviech
Ben Bethell
Umut Sagir
Arjuna James

Overfitting Overlords

Esperanza Shi
Maria Sharif
Adam Beedell
Andrew

Dropout Disco

Ethan Edwards
Dan Goss
Charles Cai
Ewan Beattie

Feature Fiestas

Tao Zamorano
Jacob Jenner
Ayman Abbas
Halil Serkan Uz

Hyperparameter Hippies

James Yan
Aparna Pillai
Clement Ha
Marcin Tolysz